

ID		Severity	Start_Time	End_Time	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	Description	...	Roundabout	Station	Stop	Traffic_Calming	Traffic_Signal
0	A-1	3	2016-02-08 00:37:08	2016-02-08 06:37:08	40.108910	-83.092860	40.112060	-83.031870	3.230	Between Sawmill Rd/Exit 20 and OH-315/Olentang...	...	False	False	False	False	False
	A-2	2	2016-02-08 05:56:20	2016-02-08 11:56:20	39.865420	-84.062800	39.865010	-84.048730	0.747	At OH-4/OH-235/Exit 41 - Accident.	...	False	False	False	False	False
	A-3	2	2016-02-08 06:15:39	2016-02-08 12:15:39	39.102660	-84.524680	39.102090	-84.523960	0.055	At I-71/US-50/Exit 1 - Accident.	...	False	False	False	False	False
	A-4	2	2016-02-08 06:51:45	2016-02-08 12:51:45	41.062130	-81.537840	41.062170	-81.535470	0.123	At Dart Ave/Exit 21 - Accident.	...	False	False	False	False	False
	A-5	3	2016-02-08 07:53:43	2016-02-08 13:53:43	39.172393	-84.492792	39.170476	-84.501798	0.500	At Mitchell Ave/Exit 6 - Accident.	...	False	False	False	False	False

```
In [73]: df.drop('ID', axis = 1, inplace = True)
df.drop('Start_Time', axis = 1, inplace = True)
df.drop('End_Time', axis = 1, inplace = True)
df.drop('End_Lat', axis = 1, inplace = True)
df.drop('End_Lng', axis = 1, inplace = True)
df.drop('Distance(mi)', axis = 1, inplace = True)
df.drop('Description', axis = 1, inplace = True)
df.drop('Number', axis = 1, inplace = True)
df.drop('Street', axis = 1, inplace = True)
df.drop('Side', axis = 1, inplace = True)
df.drop('City', axis = 1, inplace = True)
df.drop('County', axis = 1, inplace = True)
df.drop('State', axis = 1, inplace = True)
df.drop('Country', axis = 1, inplace = True)
df.drop('Timezone', axis = 1, inplace = True)
df.drop('Airport_Code', axis = 1, inplace = True)
df.drop('Weather_Timestamp', axis = 1, inplace = True)
df.drop('Wind_Direction', axis = 1, inplace = True)
df.drop('Amenity', axis = 1, inplace = True)
df.drop('Bump', axis = 1, inplace = True)
df.drop('Crossing', axis = 1, inplace = True)
df.drop('Give_Way', axis = 1, inplace = True)
df.drop('Junction', axis = 1, inplace = True)
df.drop('No_Exit', axis = 1, inplace = True)
df.drop('Railway', axis = 1, inplace = True)
df.drop('Roundabout', axis = 1, inplace = True)
df.drop('Station', axis = 1, inplace = True)
df.drop('Stop', axis = 1, inplace = True)
df.drop('Traffic_Calming', axis = 1, inplace = True)
df.drop('Traffic_Signal', axis = 1, inplace = True)
df.drop('Turning_Loop', axis = 1, inplace = True)
df.drop('Civil_Twilight', axis = 1, inplace = True)
df.drop('Nautical_Twilight', axis = 1, inplace = True)
df.drop('Astronomical_Twilight', axis = 1, inplace = True)
df.drop('Wind_Chill(F)', axis = 1, inplace = True)
df.drop('Wind_Speed(mph)', axis = 1, inplace = True)
df.drop('Precipitation(in)', axis = 1, inplace = True)
df.drop('Pressure(in)', axis = 1, inplace = True)
df.drop('Visibility(mi)', axis = 1, inplace = True)

#We are dropping variables that are either lacking in informative data (from an insurance perspective they are
# too far outside a consumer's control, or the variable themselves don't change significantly from one
# observation to the next)
```

```
In [74]: #Since we are keeping weather conditions here, we do want to get rid of the NaN weather condition values
# as they only account for 2% of the total data

df['Weather_Condition'].isna().sum() # ---> 70636, which is 2% of total weather conditions; However, also looking at
#the number of weather conditions present within that variable, it's over 130, and many of them are redundant or
#insufficiently descriptive - We will remove this now and then depending on future modeling we can decide if we want
#to also model with it included
#df.shape ---> (2845342,9)
```

Out[74]: 70636

```
In [75]: df = df[df['Weather_Condition'].notna()]
```

```
In [76]: df['Weather_Condition'].isna().sum()
```

```
#Confirmed we removed the NaN values for weather condition
```

Out[76]: 0


```
In [77]: #Before creating any dummy variables, we will start to visualize the data that we have to look for
# trends or other outstanding observations

#Severity:

df['Severity'].value_counts().plot(kind='bar')

#the overwhelming majority of accidents are Low_mid severity on a scale of Low, Low_mid, mid-high and high
```


Out[77]: <AxesSubplot:>



```
In [78]: df['Humidity(%)'].value_counts().plot(kind='bar')
```

```
#This doesn't share much with us beyond the fact that there is a typical range of humidity, but this  
# might still be helpful as we move forward
```


```
Out[78]: <AxesSubplot:>
```



```
In [79]: df['Sunrise_Sunset'].value_counts().plot(kind='bar')
```

```
#More accidents happen during the day than at night, which makes sense since more driving happens  
# during daylight hours - But is there any connection with severity and day vs night?
```

```
Out[79]: <AxesSubplot:>
```



```
In [80]: df['Temperature(F)'].value_counts().plot(kind='bar')
```

```
#Much Like humidity, temperature might just indicate that there's a certain window of temperatures  
# that lend themselves more to people being out driving
```

```
Out[80]: <AxesSubplot:>
```



```
In [81]: df.corr()
```

```
#This removes values for categorical variables, which are important for this project, though we do see  
# a lot of variation in output below, although there aren't any strong correlations
```

```
Out[81]:
```

	Severity	Start_Lat	Start_Lng	Temperature(F)	Humidity(%)
Severity	1.000000	0.090110	0.112458	-0.045177	0.037888
Start_Lat	0.090110	1.000000	-0.157846	-0.475342	0.005714
Start_Lng	0.112458	-0.157846	1.000000	0.032446	0.171060
Temperature(F)	-0.045177	-0.475342	0.032446	1.000000	-0.366342
Humidity(%)	0.037888	0.005714	0.171060	-0.366342	1.000000

In [82]: *#We can scale the data, but first we will need to set up dummy variables ...*

```
#Encoding categorical data

df = pd.concat((df, pd.get_dummies(df.Sunrise_Sunset)),1)

df.drop(['Sunrise_Sunset'], axis=1, inplace=True)
```

In [83]: df.head()

Out[83]:

	Severity	Start_Lat	Start_Lng	Zipcode	Temperature(F)	Humidity(%)	Weather_Condition	Day	Night
0	3	40.108910	-83.092860	43017	42.1	58.0	Light Rain	0	1
1	2	39.865420	-84.062800	45424	36.9	91.0	Light Rain	0	1
2	2	39.102660	-84.524680	45203	36.0	97.0	Overcast	0	1
3	2	41.062130	-81.537840	44311	39.0	55.0	Overcast	0	1
4	3	39.172393	-84.492792	45217	37.0	93.0	Light Rain	1	0

In [84]: *#Also do same for weather conditions...*

```
df = pd.concat((df, pd.get_dummies(df.Weather_Condition)), 1)
df.drop(['Weather_Condition'], axis = 1, inplace = True)
```

In [85]: *#For further standardization, I am removing anything beyond the main 5 digits of zipcodes*

```
df['Zipcode'] = df['Zipcode'].str.split('-').str[0]
```

In [86]: *#Have we gotten rid of all NaN values?*

```
#26,595 NaN values out of 2774706 observations is less than 1%, so we will delete all NaN values

df.isnull().sum().sum()
```

Out[86]: 26595

In [87]: df = df.dropna()

In [88]: df.isnull().sum().sum()

Out[88]: 0

In [89]: df2 = df

```
#We are saving a copy of df as df2 since it is cleaned except the operations being completed next
# This will serve as a placeholder in case we have issues with scaling/absolute values.
```

In [90]: *#We are now going to get absolute values on any variable that has a negative value in the df*
While scaling will fix this to an extent, it does not help when it comes time for using the
chi-squared metric for feature selection. To reduce re-work, we will proceed as below:

```
df2['Start_Lat'] = df2['Start_Lat'].abs()
df2['Start_Lng'] = df2['Start_Lng'].abs()
df2['Temperature(F)'] = df2['Temperature(F)'].abs()
df2['Humidity(%)'] = df2['Humidity(%)'].abs()
```

In [91]: df2.head()

Out[91]:

	Severity	Start_Lat	Start_Lng	Zipcode	Temperature(F)	Humidity(%)	Day	Night	Blowing Dust	Blowing Dust / Windy	...	Thunder and Hail / Windy	Thunder in the Vicinity	Thunderstorm	Thunderstorms and Rain	Tornado	Volcanic Ash
0	3	40.108910	83.092860	43017	42.1	58.0	0	1	0	0	...	0	0	0	0	0	0
1	2	39.865420	84.062800	45424	36.9	91.0	0	1	0	0	...	0	0	0	0	0	0
2	2	39.102660	84.524680	45203	36.0	97.0	0	1	0	0	...	0	0	0	0	0	0
3	2	41.062130	81.537840	44311	39.0	55.0	0	1	0	0	...	0	0	0	0	0	0
4	3	39.172393	84.492792	45217	37.0	93.0	1	0	0	0	...	0	0	0	0	0	0

5 rows × 135 columns

In [92]: *#Splitting our data into training and test sets -*

```
x = df2.loc[ : , df2.columns != 'Severity'] # Features
y = df2.Severity # Target variable

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1) # 80% training and 20% test
```

```

In [93]: #Test Set

xt = X_test.values
xt.shape

#feature matrix

Out[93]: (551905, 134)

In [94]: #Train Set

xs = X_train.values
xs.shape

#feature matrix

Out[94]: (2207618, 134)

In [95]: xr = df2.values
xr.shape

#feature matrix for full dataset of df2

Out[95]: (2759523, 135)

In [96]: #Now we will scale the data -

scaler = StandardScaler()

In [97]: scaler.fit(xs)

Out[97]: StandardScaler()

In [98]: xs_scaled = scaler.transform(xs)

#Training values scaled

In [99]: pca = PCA(n_components = 0.99, whiten = True)

#Using PCA to investigate relevant features

In [100]: features_pca = pca.fit_transform(xs)

In [101]: features_pca.shape

#This shows a reduction from 135 to 1 when done on our training set
#We can also check for its impact on the test and the full sets...

Out[101]: (2207618, 1)

In [102]: #Test fit

scaler.fit(xt)

Out[102]: StandardScaler()

In [103]: features_pca_test = pca.fit_transform(xt)

In [104]: features_pca_test.shape

#This reduced features significantly - 135 down to 1

Out[104]: (551905, 1)

In [105]: #For the full set

scaler.fit(xr)

Out[105]: StandardScaler()

In [106]: xr_scaled = scaler.transform(xr)

In [107]: features_pca_full = pca.fit_transform(xr)

In [108]: features_pca_full.shape

#Same reduction compared to the original - down to 1 from 135

Out[108]: (2759523, 1)

In [109]: '''

Using transformed PCA data to again get R^2 score and RMSE

'''

Out[109]: '\n\nUsing transformed PCA data to again get R^2 score and RMSE\n\n'

```

```
In [110]: regression = linear_model.LinearRegression()

#While I am using a Decision Tree Model, as I build up to that, I am also exploring any Linear relationship
# that could exist in the data - We will quickly see there is not
```

```
In [111]: regression.fit(features_pca, y_train)

#Fitting the model
```

```
Out[111]: LinearRegression()
```

```
In [112]: y_c = regression.predict(features_pca)

#Using the model for predictions
```

```
In [113]: y_cv = cross_val_predict(regression, features_pca, y_train, cv = 10)

#Cross validating
```

```
In [114]: score_c = r2_score(y_train, y_c)
score_cv = r2_score(y_train, y_cv)

#Obtaining evaluation metrics - R^2
```

```
In [115]: mse_c = mean_squared_error(y_train, y_c)
mse_cv = mean_squared_error(y_train, y_cv)

#Obtaining evaluation metrics - MSE
```

```
In [116]: score_c

# 0.8 % - R^2 score
# This should indicate our model is not a good fit - A Linear Model is not our best option
```

```
Out[116]: 0.008610448812355287
```

```
In [117]: mse_c

#MSE
```

```
Out[117]: 0.22602142266055894
```

```
In [118]: '''
At this point, we will move forward with the Decision Tree Model
'''
```

```
Out[118]: '\nAt this point, we will move forward with the Decision Tree Model\n'
```

```
In [119]: # Create Decision Tree classifier object
clf = DecisionTreeClassifier()

# Train Decision Tree Classifier
clf = clf.fit(X_train,y_train)

# Predict the response for test dataset
y_pred = clf.predict(X_test)
```

```
In [120]: #Now we will report the accuracy of the model and create a confusion matrix for
# the model prediction on the test set

y_pred = clf.predict(X_test)

print("Model Accuracy with criterion gini index of {0:0.4f}". format(accuracy_score(y_test, y_pred)))

print(confusion_matrix(y_test, y_pred))

#Accuracy is 87% which seems a bit high, but perhaps due to already disproportionately high
# instances of Losses of 2nd Level Severity?

Model Accuracy with criterion gini index of 0.8723
[[ 2222  2523   231   130]
 [ 2975 457978 17815 12819]
 [   273 16841 10622  2317]
 [   137 12299   2097 10626]]
```

In [121]: *#We can also optimize this tree:*

```
# Create Decision Tree classifier object
clf = DecisionTreeClassifier(criterion="entropy", max_depth=3)

# Train Decision Tree Classifier
clf = clf.fit(X_train,y_train)

#Predict the response for test data
y_pred = clf.predict(X_test)

# Model Accuracy
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))

#We see here 89% which is a slight improvement over our initial iteration of the model
```

Accuracy: 0.8907094518078292

In [122]: *# Making Predictions with Our Model*

```
predictions = clf.predict(X_test)
print(predictions[:5])

#5 predictions indicating the predicted severity of 5 losses
# We see all 2s, which may go back to the original concern that there are such a high proportion of
# Losses rated a 2 in severity that a default assumption of 2 Level severity is as good a guess as
# what the model shows
```

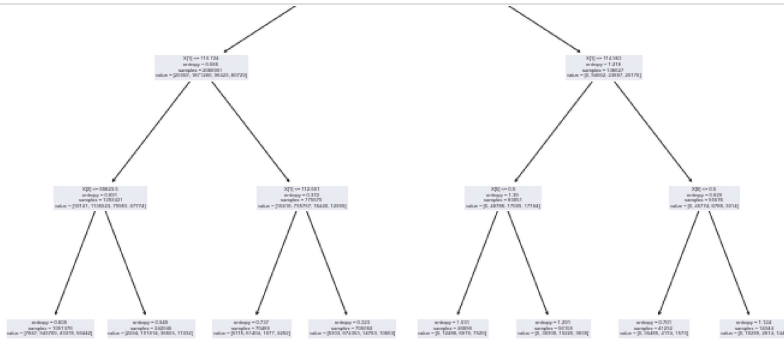
[2 2 2 2 2]

In [123]: *#We will now plot and visualize the Decision Tree*

```
plt.figure(figsize = (12, 8))

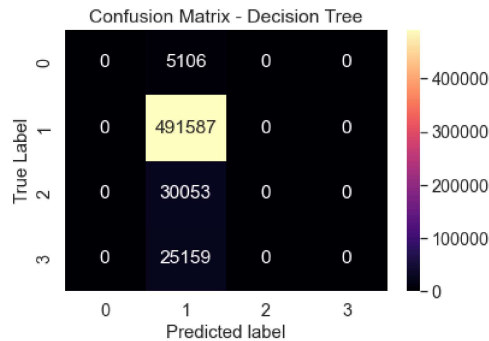
from sklearn import tree

tree.plot_tree(clf.fit(X_train, y_train))
```



In [124]: *#Visualizing with Confusion Matrix*

```
#import the relevant packages
from sklearn import metrics
import seaborn as sns
import matplotlib.pyplot as plt#get the confusion matrix
confusion_matrix = metrics.confusion_matrix(y_test,
                                             y_pred)#turn this into a dataframe
matrix_df = pd.DataFrame(confusion_matrix)#plot the result
ax = plt.axes()
sns.set(font_scale=1.3)
plt.figure(figsize=(10,7))
sns.heatmap(matrix_df, annot=True, fmt="g", ax=ax, cmap="magma")#set axis titles
ax.set_title('Confusion Matrix - Decision Tree')
ax.set_xlabel("Predicted label", fontsize =15)
ax.set_ylabel("True Label", fontsize=15)
plt.show()
```



In [125]: *#Now we can use the Chi-Squared metric to try to select the top 5 features:*

```
sel5 = SelectKBest(score_func = chi2, k = 5)
sel5.fit(df2.fillna(0), y)
df2.columns[sel5.get_support()].to_numpy()
```

Out[125]: array(['Severity', 'Start_Lng', 'Zipcode', 'Humidity(%)', 'Clear'],
dtype=object)

In [126]: *#Another feature selection option:*

```
#Create a feature list from our dataframe
feature_list = list(df2.columns)

#Getting numerical feature importance:
importances = list(clf.feature_importances_)

# List of tuples with variables and importance
feature_importances = [(feature, round(importance, 2)) for feature, importance in zip(feature_list, importances)]

# Sorting the feature importance in descending order by importance
feature_importances = sorted(feature_importances, key = lambda x: x[1], reverse = True)
# Printing the feature and importances
[print('Variable: {:20} Importance: {}'.format(*pair)) for pair in feature_importances];
```

*#Note - This is different from the previous List of top features, and this actually
accounts for the fact that Severity itself should be removed from the calculation.
#This helps us see that Latitude, snow, Longitude, humidity and day vs. night do play a part*

```
Variable: Severity           Importance: 0.0
Variable: Zipcode            Importance: 0.0
Variable: Temperature(F)    Importance: 0.0
Variable: Night              Importance: 0.0
Variable: Blowing Dust       Importance: 0.0
Variable: Blowing Dust / Windy Importance: 0.0
Variable: Blowing Sand       Importance: 0.0
Variable: Blowing Snow       Importance: 0.0
Variable: Blowing Snow / Windy Importance: 0.0
Variable: Clear              Importance: 0.0
Variable: Cloudy              Importance: 0.0
Variable: Cloudy / Windy     Importance: 0.0
Variable: Drifting Snow      Importance: 0.0
Variable: Drizzle            Importance: 0.0
Variable: Drizzle / Windy    Importance: 0.0
Variable: Drizzle and Fog    Importance: 0.0
Variable: Dust Whirls        Importance: 0.0
Variable: Duststorm          Importance: 0.0
Variable: Fair                Importance: 0.0
Variable: Fair / Windy       Importance: 0.0
```



```
In [127]: #Now we will use these top 5 features to do what we did above -
# We will fit a decision tree classifier on training set and then report the accuracy
# and create a confusion matrix for the model prediction on test sets

top5 = sel5.transform(df2)
top5 = pd.DataFrame(top5)
```

```
In [128]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(top5, y, test_size = 0.2)

clf = DecisionTreeClassifier()

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)
print("Model Accuracy with index of {0:0.4f}".format(accuracy_score(y_test, y_pred)))

from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred))

#Here we see that when we limit to 5 variables we get an index score of 1.0 which means
# that there is no relationship between the variables - that is the results are possibly
# totally random using the model as-is

Model Accuracy with index of 1.0000
[[ 5160    0    0    0]
 [    0 491672    0    0]
 [    0    0 30131    0]
 [    0    0    0 24942]]
```

In []:

In []:

In []:

```
In [129]: '''
Another model I explored was a totally different approach that looked at the text descriptions of the
accidents to see if any helpful information could be gleaned from that data
'''
```

```
Out[129]: '\nAnother model I explored was a totally different approach that looked at the text descriptions of the\naccidents to see if any helpful infor
mation could be gleaned from that data\n'
```

```
In [130]: #Importing Libraries

import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.feature_extraction.text import TfidfVectorizer
from scipy.sparse import hstack
from matplotlib import pyplot as plt
import seaborn as sns
import eli5
```

```
In [131]: #Loading data

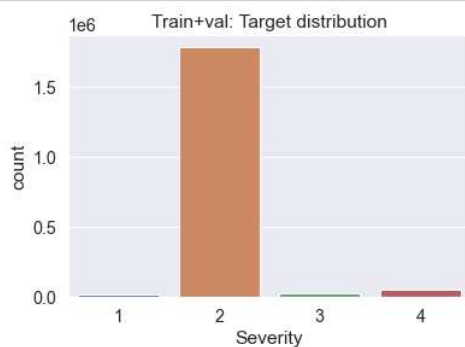
train = pd.read_csv('US_Accidents_Dec21_updated.csv', index_col='ID').dropna()
valid = pd.read_csv('US_Accidents_Dec21_updated.csv', index_col='ID').dropna()
test = pd.read_csv('US_Accidents_Dec21_updated.csv', index_col='ID').dropna()
```

```
In [132]: #Concatenating training and validation sets

train_val = pd.concat([train, valid])
```

```
In [133]: #Finding distribution of Loss severity (our target) across training and validation data sets

sns.countplot(train_val['Severity']);
plt.title('Train+val: Target distribution');
```



```
In [134]: #Setting up text transformer
```

```
text_transformer = TfidfVectorizer(stop_words='english', ngram_range=(1, 2), lowercase=True, max_features=150000)
```

```
In [136]: #Transforming training and testing text
```

```
X_train_text = text_transformer.fit_transform(train_val['Description'])  
X_test_text = text_transformer.transform(test['Description'])
```

```
In [137]: #Data dimensions
```

```
X_train_text.shape, X_test_text.shape
```

```
Out[137]: ((1886636, 150000), (943318, 150000))
```

```
In [138]: #Log Reg for our text data
```

```
logit = LogisticRegression(C=5e1, solver='lbfgs', multi_class='multinomial', random_state=17, n_jobs=4)
```

```
In [139]: #Cross Validation
```

```
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=17)
```

```
In [140]: cv_results = cross_val_score(logit, X_train_text, train_val['Severity'], cv=skf, scoring='f1_micro')
```

```
In [141]: cv_results, cv_results.mean()
```

```
Out[141]: (array([0.96953579, 0.96909842, 0.96921238, 0.96954101, 0.96922298]),  
          0.9693221160766188)
```

```
In [142]: #Fitting the model
```

```
logit.fit(X_train_text, train_val['Severity'])
```

```
Out[142]: LogisticRegression(C=50.0, multi_class='multinomial', n_jobs=4, random_state=17)
```

```
In [143]: #Finding top text features

eli5.show_weights(estimator=logit,
                  feature_names= list(text_transformer.get_feature_names()),
                  top=(50, 5))
```

Out[143]:

y=1 top features		y=2 top features		y=3 top features		y=4 top features	
Weight'	Feature	Weight'	Feature	Weight'	Feature	Weight'	Feature
+8.733	road accident	+39.874	traffic	+8.010	accident	+42.264	closed
+8.662	rd accident	+29.886	stationary traffic	+6.329	lane closed	+20.873	closed accident
+7.662	ave accident	+29.886	stationary	+5.776	il	+15.983	closed road
+6.674	60 grand	+26.114	slow traffic	+5.028	belt	+11.668	road closed
+6.038	accident lanes	+26.066	slow	+4.159	fullerton ave	+10.790	blocked
+5.559	st accident	+23.379	incident	+4.002	accident lanes	+8.117	rd
+5.389	earlier accident	+21.533	near	+3.885	rd accident	+6.554	expect delays
+5.334	83rd ave	+20.670	caution	+3.810	dr accident	+6.085	road
+5.183	drexel rd	+19.237	drive caution	+3.773	cedar lake	+5.590	expect
+4.970	earlier	+17.695	drive	+3.735	traffic problem	+5.520	closed cr
+4.880	az	+14.965	crash	+3.688	hwy accident	+5.249	closure
+4.817	magee rd	+10.724	shoulder closed	+3.650	county farm	+5.135	delays
+4.781	dr accident	+10.430	fl	+3.643	pulaski	+5.058	mp expect
+4.714	magee	+10.303	veh	+3.611	lane traffic	+4.668	lanes blocked
+4.654	kolb rd	+10.061	traffic fl	+3.551	army	+4.591	lane blocked
+4.645	kolb	+9.825	alternate route	+3.489	accident traffic	+4.536	crash investigation
+4.627	thornsdale	+9.136	conndot	+3.449	gary ave	+4.474	milemarker
+4.627	thornsdale rd	+8.920	closed alternate	+3.417	overturned	+4.438	fl
+4.390	rd earlier	+8.618	right shoulder	+3.373	overturned vehicle	+4.395	summit rd
+4.350	tn	+8.437	shoulder	+3.370	20 lake	+4.141	county
+4.293	sw	+8.256	use	+3.349	sherman dr	+4.010	investigation
+4.288	tangerine rd	+7.918	1039	+3.347	kedzie	+3.992	ny
+4.285	craycroft rd	+7.818	accident	+3.315	pulaski rd	+3.647	summit
+4.275	craycroft	+7.796	alternate	+3.282	lane	+3.632	blocked overturned
+4.264	drexel	+7.630	road	+3.261	accident lane	+3.593	95 accident
+4.259	blvd accident	+7.472	lanes closed	+3.233	army trail	+3.520	motorists
+4.201	59th ave	+7.446	right	+3.223	lanes blocked	+3.514	st
+4.161	tatum blvd	+7.167	lanes	+3.217	alternate lane	+3.416	twp lanes
+4.160	mcdowell rd	+7.132	rd drive	+3.213	single alternate	+3.249	nb near
+4.097	tangerine	+6.943	ave	+3.121	bartlett rd	+3.172	126 closed
+4.094	prince rd	+6.904	route	+3.104	accident left	+3.158	directions
+4.070	ne	+6.778	lane closed	+3.052	ryan	+3.127	rt
+3.918	grant rd	+6.546	rd	+3.041	blocked right	+3.076	construction ca
+3.877	pike accident	+6.391	mn	+3.038	st accident	+3.043	route
+3.861	speedway blvd	+6.357	directions road	+3.033	single	+3.042	closed fl
+3.847	35th ave	+6.173	closed incident	+3.014	mcculloch blvd	+3.024	sb near
+3.838	cave creek	+6.074	use alternate	+2.961	ave exit	+3.018	nj
+3.813	dunlap ave	+5.968	lane lanes	+2.957	problem	+2.962	126
+3.800	accident	+5.903	chp	+2.941	blocked ahead	+2.959	95
+3.737	91st ave	+5.893	rd near	+2.928	purcell blvd	+2.915	motorists expect
+3.719	litchfield rd	+5.749	ca	+2.925	butterfield rd	+2.881	traffic affected
+3.711	swan rd	+5.651	accident road	+2.890	riverwoods rd	+2.801	road 126
+3.708	tatum	+5.571	spun	+2.886	ahead	+2.796	eb near
+3.687	valencia rd	+5.534	vehicle spun	+2.878	midlothian rd	+2.765	dr
+3.684	mcclintock dr	+5.273	st drive	+2.859	exit	+2.751	22 closed
+3.683	pima	+5.262	vs	+2.828	lombard accident	+2.748	st directions
+3.656	harrison rd	+5.231	lane	+2.824	tollway	+2.746	wb near
+3.634	litchfield	+5.191	restriction	+2.805	accident blocked	+2.705	rd road
+3.633	mcclintock	+5.160	near house	+2.786	roosevelt rd	+2.645	mn
+3.621	rural rd	+4.998	tc	+2.776	gary	+2.637	md
... 6742 more positive 119927 more positive 11223 more positive 25905 more positive ...	
... 143204 more negative 30019 more negative 138723 more negative 124041 more negative ...	
-9.903	near	-6.290	accident right	-10.837	slow	-8.239	conndot
-10.768	stationary traffic	-6.622	blocked	-11.449	closed	-8.492	rd accident
-10.768	stationary	-12.608	closed road	-12.502	stationary	-8.999	shoulder closed
-10.774	rd	-17.022	closed accident	-12.502	stationary traffic	-15.620	lane closed
-16.535	traffic	-24.230	closed	-15.728	traffic	-19.628	accident

```
In [144]: test_preds = logit.predict(X_test_text)

In [145]: pd.DataFrame(test_preds, columns=['Severity']).head()
```

Out[145]:

	Severity
0	4
1	4
2	4
3	2
4	2

```
In [ ]:
```